BIG DATA
DEVELOPMENT

ACAD**GILD**

Session 3: HDFS Internals

Assignment 1

Assignment 1 – Try the given quiz questions and provide the answers in a word document.

Table of Contents

Big Data and Hadoop Development

1. Introduction

In this assignment you need to select one right choice for the questions given on the topics discussed in the third session.

2. Objective

This assignment will help you to consolidate the concepts learnt in the session 3.

3. Prerequisites:
None

4. Associated Data Files
None

5. Problem Statement:

1. HDFS is built around the idea that data is written _____but read many times.
a) many
b) twice
c) data already exists
d) once

2. Hadoop divides input into fixed size pieces called what?
a) output result
b) input splits
c) input data
d) input blogs

3. All the blocks are replicated in other nodes for _____
a) security
b) big data
c) pool
d) fault tolerance

4. Block size can be changed using the properties in _____

a) core-site.xml

b) Hadoop-env.sh

c) hdfs-site.xml

d) yarn-site.xml

5. Hadoop uses the _____representation of the data stored in the file blocks known as Input splits.

a) physical

b) logical

c) mechanical

d) none

6. DFS calls NameNode to create file in file system's_____

a) dataspace

b) resourcespace

c) namespace

d) nodespace

7. Data packets are streamed to first DataNode in the _____

a) handshake

b) pipeline

c) hard disk

d) hdfs

8. The client has finished writing data, it calls _____on the stream.

a) close()

b) read()

c) open()

d) check()

9. Blocks are read in order, with the _____ opening new connections to datanodes as the client reads through the stream.

a) DFSoutputstream
b) <mark>DFSInputStream</mark>
c) DFStrackManager
d) DFSStringConcatination

10. If I have 100 input splits, how many maps will run?

a) 200
b) 50
c) <mark>100</mark>
d) 1

6. Expected Output

None

7. Approximate Time to Complete Task

15 mins